SuperSpike BPTT reference saturates: 1.0000 ff / 1.0000 rec, every suite, both graphs. Every other rule tested clears the gate, so this panel shows WHICH SINGLE RULE FAILS a task the rest saturate; it does not rank the rest. Pass / fail / at-chance, not a graded surface.

	Low addressability (broadcast)	High addressability (directed / local feedback)
Low richness (±1)	**FAIL** ±1 × surrogate eligibility MatchedLocal: the lead FAIL, at chance 0.5000 ff / 0.5100 rec **contrast (not gated)** ±1 broadcast REINFORCE MatchedRlFlat: same reward and topology 0.7775 ff / 0.7962 rec One cell, two rules, 0.28 apart. Collapsing them into one “broadcast ±1” box is the overreach the lead claim’s wording exists to avoid.	**PASS** REINFORCE × frozen B_i MatchedRlReinforceFb: directed, same ±1 reward 0.9950 ff / 0.9812 rec Richness alone does not decide it: a graded BROADCAST rule reaches 0.9975. Addressability alone does not either: ±1 through a frozen B_i passes. What fails is the one rule that has neither.
High richness (graded)	**contrast (not gated)** broadcast graded error MatchedBroadcastGradedError: not a locality proof 0.9975 ff / 0.9975 rec	**PASS** graded error × DFA MatchedDfaGradedError 0.9925 ff / 0.9875 rec

Panel B: XOR locality flip (1-layer xor_thresh). Addressability evidence: broadcast fails a task DFA solves.

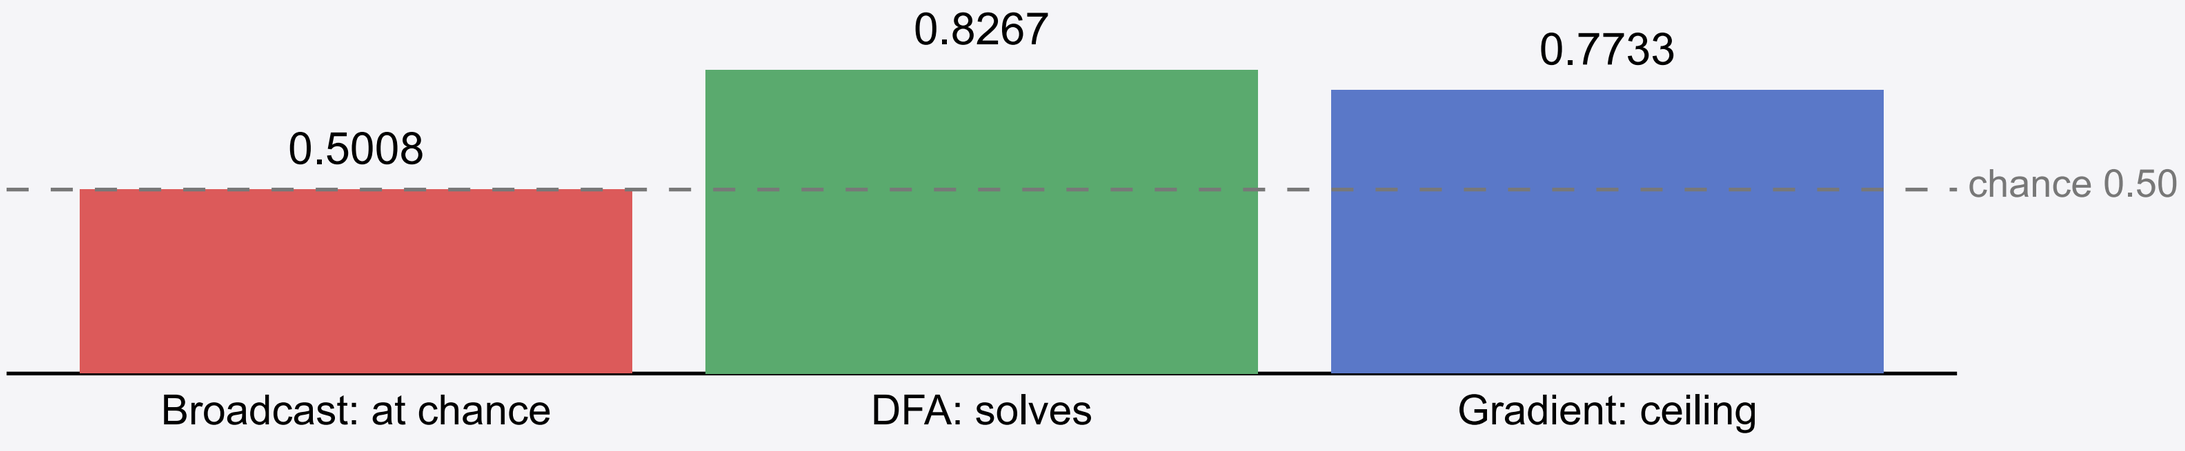

Not claimed for 2-layer mid-init depth locality, where broadcast also solves. Matched PASS does not imply live muted-θ / k-WTA G2 PASS.