h128 / published-2ms / e400 / d32-L4, seed-paired. Does the read-out ADD temporal structure no such substrate represents, or SUBSTITUTE for the adaptation and recurrence ff+fixed happens not to have?

Panel A: one shared accuracy axis, in the row order of Table SHD-7. NOT sorted by gain.

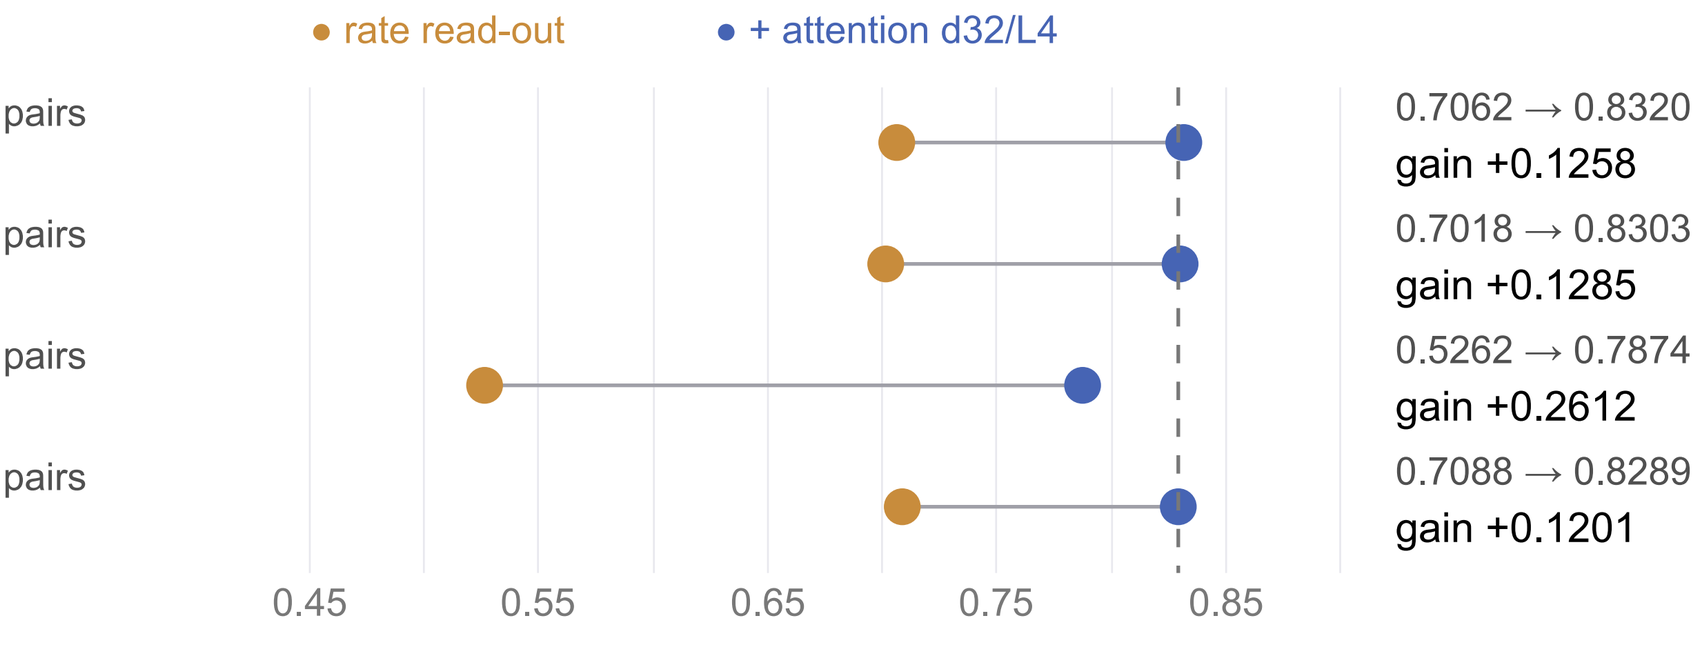

ff+fixed the anchor	scale 1.0 · n = 12 pairs	0.7062 → 0.8320 gain +0.1258
ff+alif threshold adaptation	scale 1.0 · n = 12 pairs	0.7018 → 0.8303 gain +0.1285
rec+alif recurrent + adaptation	scale 0.4 · n = 10 pairs	0.5262 → 0.7874 gain +0.2612
ff+fixed scale-matched control	scale 0.4 · n = 12 pairs	0.7088 → 0.8289 gain +0.1201

THE RECURRENT SUBSTRATE DOES NOT WIN. rec+alif+attn reaches 0.7874 against ff+fixed+attn's 0.8289 at the same scale: attention closes most of the gap the substrate gives away, and not all of it. NO VERDICT IS ISSUED ON THAT ORDERING, and nothing here is sorted by gain.

Panel B: the registered contrasts, each against its two-sided bar

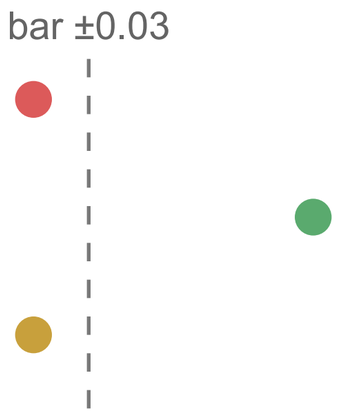

A-1 gain(ff+alif) − gain(ff+fixed)	+0.0027 6 of 12	a coin flip: adaptation makes no difference to the gain HERE
M-2 gain(rec+alif) − gain(ff+fixed), scale 0.4	+0.1411 10 of 10	SUPPORTED: the gain roughly doubles on the recurrent substrate
M-4 ff+fixed at scale 0.4 vs archived 1.0	+0.0026 —	the scale is not doing the work

BOTH READINGS OF “THE GAIN DOUBLES”, because neither may be quoted alone. rec+alif starts 0.18 LOWER, with 0.4738 of headroom against 0.2912.

Gain / headroom is 0.551 against 0.412, so the ratio falls 2.2× → 1.34×. The ORDERING SURVIVES; most of its apparent size does not. That normalisation is POST-HOC AND NOT REGISTERED, and saying so is not licence to prefer whichever of the two numbers is more convenient, in either direction.

Panel C: usability, which had to be measured before any row of Panel A could be. This is why the recurrent row exists at one scale only.

arm	scale	completed	voided	diverged
rec+alif	0.4	11 / 12	0	1
rec+alif	1.0	8 / 12	0	4
rec+fixed	0.4	7 / 12	5	0
rec+fixed	1.0	5 / 12	5	2

rec+fixed's ten voids are ALL SATURATION (up to 52.3% of hidden units pinned at maximum firing), and NONE is divergence at scale 0.4. So on the recurrent substrate ADAPTATION IS STABILISING, the opposite of the sign that wave's own hypothesis asserted.

A-1 above is therefore inert AT THIS OPERATING POINT, on this substrate, not in general.

TEN PAIRS (10 of 12), THE REGISTERED MINIMUM. The two recurrent arms lost different seeds and one further loss on either would have made M-2 unreportable. Pairing REDUCES survivorship rather than removing it: the surviving pairs are those that did not diverge, and divergence is not random. The feed-forward rows carry no such exposure at 12/12.