h128 / published-2ms / e400 / d32-L4, seed-paired. Does the read-out ADD temporal structure no such substrate represents, or SUBSTITUTE for the adaptation and recurrence ff+fixed happens not to have?

Panel A — one shared accuracy axis, in the row order of Table SHD-7. NOT sorted by gain.

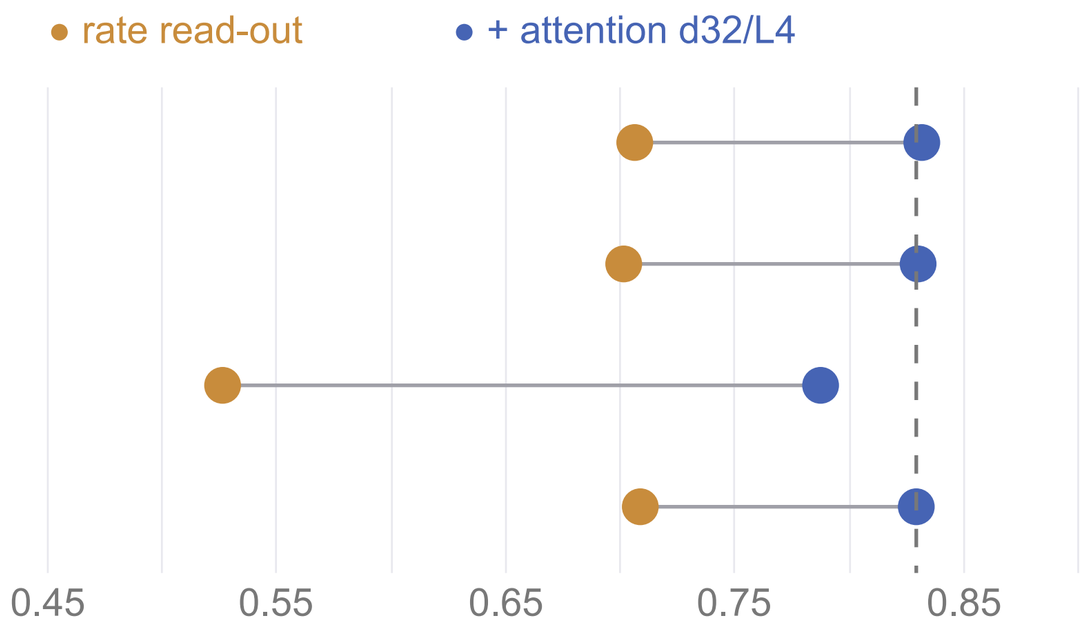

ff+fixed — the anchor	scale 1.0 · n = 12 pairs	0.7062 → 0.8320 gain +0.1258
ff+alif — threshold adaptation	scale 1.0 · n = 12 pairs	0.7018 → 0.8303 gain +0.1285
rec+alif — recurrent + adaptation	scale 0.4 · n = 10 pairs	0.5262 → 0.7874 gain +0.2612
ff+fixed — scale-matched control	scale 0.4 · n = 12 pairs	0.7088 → 0.8289 gain +0.1201

THE RECURRENT SUBSTRATE DOES NOT WIN. rec+alif+attn reaches 0.7874 against ff+fixed+attn's 0.8289 at the same scale: attention closes most of the gap the substrate gives away, and not all of it. NO VERDICT IS ISSUED ON THAT ORDERING, and nothing here is sorted by gain.

Panel B — the registered contrasts, each against its two-sided bar

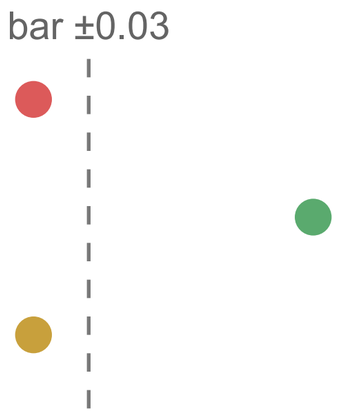

A-1 gain(ff+alif) − gain(ff+fixed)	+0.0027 6 of 12	a coin flip: adaptation makes no difference to the gain HERE
M-2 gain(rec+alif) − gain(ff+fixed), scale 0.4	+0.1411 10 of 10	SUPPORTED: the gain roughly doubles on the recurrent substrate
M-4 ff+fixed at scale 0.4 vs archived 1.0	+0.0026 —	the scale is not doing the work

BOTH READINGS OF “THE GAIN DOUBLES”, because neither may be quoted alone. rec+alif starts 0.18 LOWER, with 0.4738 of headroom against 0.2912.

Gain / headroom is 0.551 against 0.412, so the ratio falls 2.2× → 1.34×. The ORDERING SURVIVES; most of its apparent size does not. That normalisation is POST-HOC AND NOT REGISTERED, and saying so is not licence to prefer whichever of the two numbers is more convenient, in either direction.

Panel C — usability, which had to be measured before any row of Panel A could be. This is why the recurrent row exists at one scale only.

arm	scale	completed	voided	diverged
rec+alif	0.4	11 / 12	0	1
rec+alif	1.0	8 / 12	0	4
rec+fixed	0.4	7 / 12	5	0
rec+fixed	1.0	5 / 12	5	2

rec+fixed's ten voids are ALL SATURATION — up to 52.3% of hidden units pinned at maximum firing — and NONE is divergence at scale 0.4. So on the recurrent substrate ADAPTATION IS STABILISING, the opposite of the sign that wave's own hypothesis asserted.

A-1 above is therefore inert AT THIS OPERATING POINT, on this substrate — not in general.

TEN PAIRS (10 of 12), THE REGISTERED MINIMUM. The two recurrent arms lost different seeds and one further loss on either would have made M-2 unreportable. Pairing REDUCES survivorship rather than removing it: the surviving pairs are those that did not diverge, and divergence is not random. The feed-forward rows carry no such exposure at 12/12.