


What a time-axis read-out buys is temporal order

A difference-in-differences on the GAIN, on SHD. Seed-paired, n = 32 at the anchor, h128 / published-2ms / d32-L4 / e400.

NOT SHOWN HERE, AND NOT CLAIMED: that SHD depends on temporal order. That is established — Cramer et al. 2022 could not exceed 60% on spike-count-only SHD; the Neuromorphic Sequential Arena removes temporal processing model-side and reports 86.48 → 68.51; Yu et al. 2025 randomise spike times at fixed counts.

Three independent destruction operators, one conclusion, all of it prior to this work. What is measured here is WHICH COMPONENT'S contribution is the order-dependent one.



THE EFFECT'S SIZE IS NOT THE GAIN. Spearman ρ between the six per-width gains and their difference-in-differences is -0.1430, against a registered bar of +0.829.

h768 carries the SMALLEST positive gain on the ladder (+0.0560) and the LARGEST difference-in-differences in the campaign (+0.1881).

The claim is that the read-out's contribution IS order-dependent. It is NOT that the gain decomposes into an order-dependent share and a remainder.

Accuracy, for scale and no more: the read-out reaches 0.8332 against the rate arm's 0.7057. It is NOT COMPETITIVE — the SHD frontier is 95–96.4% via learned delays, adaptation and spiking transformers, and this instrument carries no temporal kernel of any kind. No accuracy claim is made and none is drawn larger than the costs above.

SCOPE, on the image because an abstract travels alone: the gain INVERTS at width h1024 (-0.1618 at d32/L4), the collapse there is LOCATED BUT UNEXPLAINED, and 0.80 clearance is geometry-specific. Three preregistered rescue levers at h1024 all failed and every one is worse than the arm it was meant to rescue.

No biology, no cortex, no neuromorphic-hardware claim, and no claim that attention makes a network temporal.