Panel A — this instrument, seed-paired. THIS IS NOT COMPETITIVE, and the axis below says so.

n = 32 (confirmation): 0.7057 → 0.8332, gain +0.1275, positive 32/32, ≥ 0.80 in 32/32, |e400 − e200| = 0.0002

n = 12 (registered): 0.7062 → 0.8320, gain +0.1258, positive 12/12, ≥ 0.80 in 12/12. Neither supersedes the other.

Panel B — the published field, on the same axis. The instrument carries no temporal kernel of any kind.

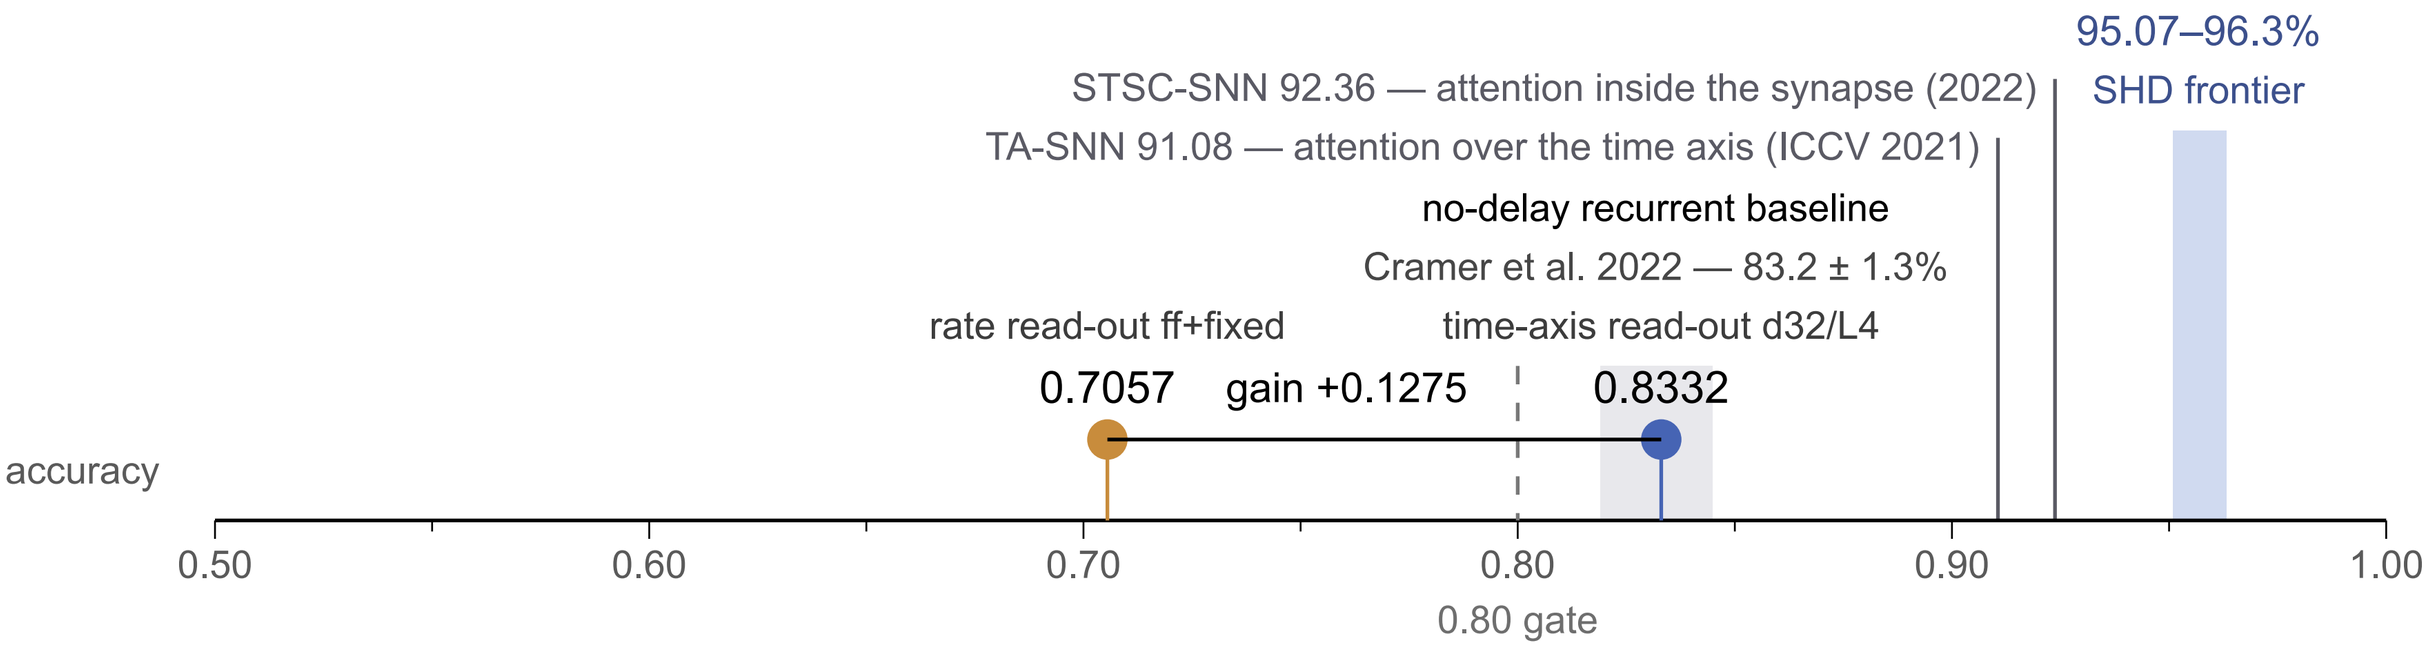

1.5 points, to scale. SHD ships no validation set: Baronig et al. 2025 report one model at 95.81 validating on test and 93.79 on a held-out split. Differences narrower than this between published SHD numbers are not reliably meaningful — including this paper's own.

Panel C — the 0.80 clearance is geometry-specific. The gain is not: attention buys roughly the same amount everywhere tested.

adjacent-sum-5 / published-2ms (anchor)	rate 0.7062	attention 0.8320	gain +0.1258	≥ 0.80 in 12/12
published-10ms	rate 0.6734	attention 0.8225	gain +0.1491	≥ 0.80 in 10/12
channels-700 (standard 700-channel input)	rate 0.6774	attention 0.7864	gain +0.1090	≥ 0.80 in 6/12

This instrument lands INSIDE the no-delay recurrent band — which is where an architecture carrying no temporal kernel of any kind should land, and is the whole reading.

Literature values in Panel B are published figures, checked against their primary sources on 2026-09-03, before submission. They are NOT MACHINE-CHECKED against cells: no cell on disk produced them.