

Figure 2 — Headline accuracy, and the 0.80 clearance

The read-out takes the instrument from 0.7057 to 0.8332, every seed positive and every seed at or above 0.80. THIS IS NOT COMPETITIVE, and the axis below says so.

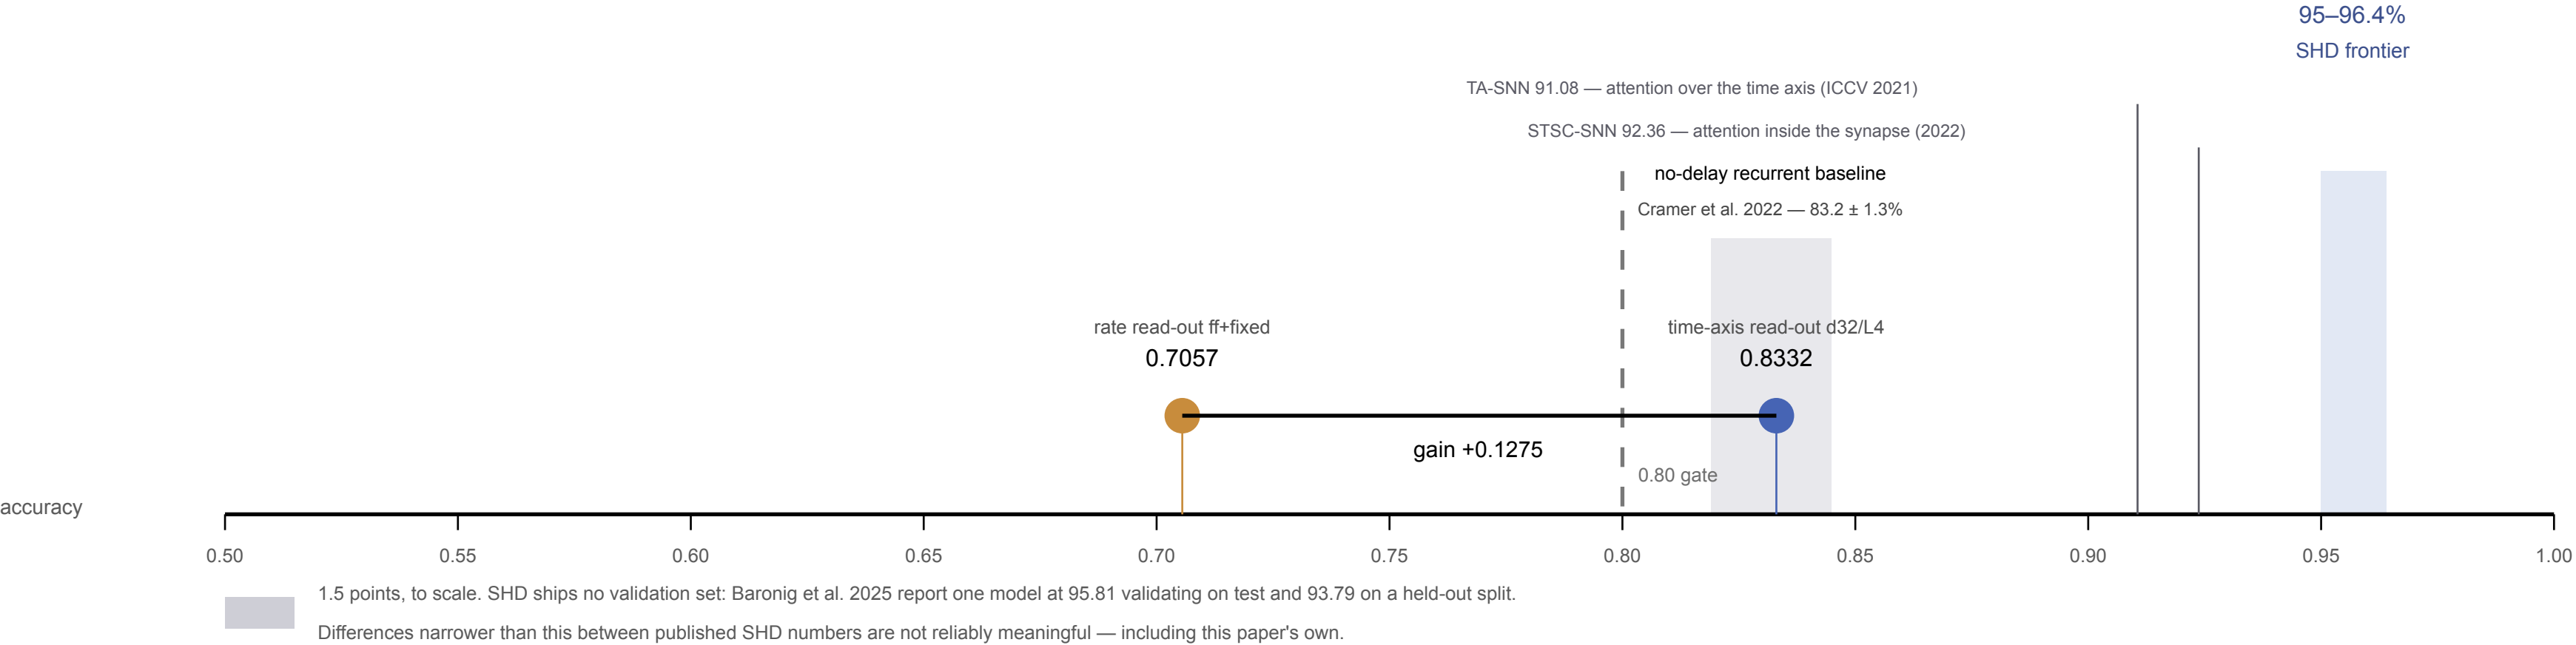
Panel A — this instrument, n = 32 seed-paired

n = 32 (confirmation): 0.7057 → 0.8332, gain +0.1275, positive 32/32, ≥ 0.80 in 32/32, |e400 – e200| = 0.0002

n = 12 (registered): 0.7062 → 0.8320, gain +0.1258, positive 12/12, ≥ 0.80 in 12/12

n = 12 is the registered measurement and n = 32 the confirmation. Neither supersedes the other.

Panel B — the published field, on the same axis. The instrument carries no temporal kernel of any kind.



Panel C — the 0.80 clearance is geometry-specific. The gain is not: attention buys roughly the same amount everywhere tested.

adjacent-sum-5 / published-2ms (anchor)	rate 0.7062	attention 0.8320	gain +0.1258	≥ 0.80 in 12/12
published-10ms	rate 0.6734	attention 0.8225	gain +0.1491	≥ 0.80 in 10/12
channels-700 (standard 700-channel input)	rate 0.6774	attention 0.7864	gain +0.1090	≥ 0.80 in 6/12

This instrument lands INSIDE the no-delay recurrent band — which is where an architecture carrying no temporal kernel of any kind should land, and is the whole reading.

Excluded from this axis on purpose: Pfa-SNN 96.26, Event-SSMA 95.90, SpikeSCR 95.60, d-cAdLIF 94.85 — a secondary comparison table, not a primary source.

NOT MACHINE-CHECKED: every literature value in Panel B comes from a 2026-08-27 search pass. Unlike every SHD number in this paper it is not verified against cells on disk, and check_every_number.py does not sweep the section it lives in. Each must be checked against its primary source before submission.