Panel A: attention's gain on six widths. n = 12 per rung, e400, seed-paired. Rungs are points: no curve is fitted and none may be.

EVERY RUNG IS d32/L4. The gain inverts at h1024 as a property of THIS read-out depth, not of the width: at h1024, d32/L2 gains +0.0405 in 20/20 seeds.

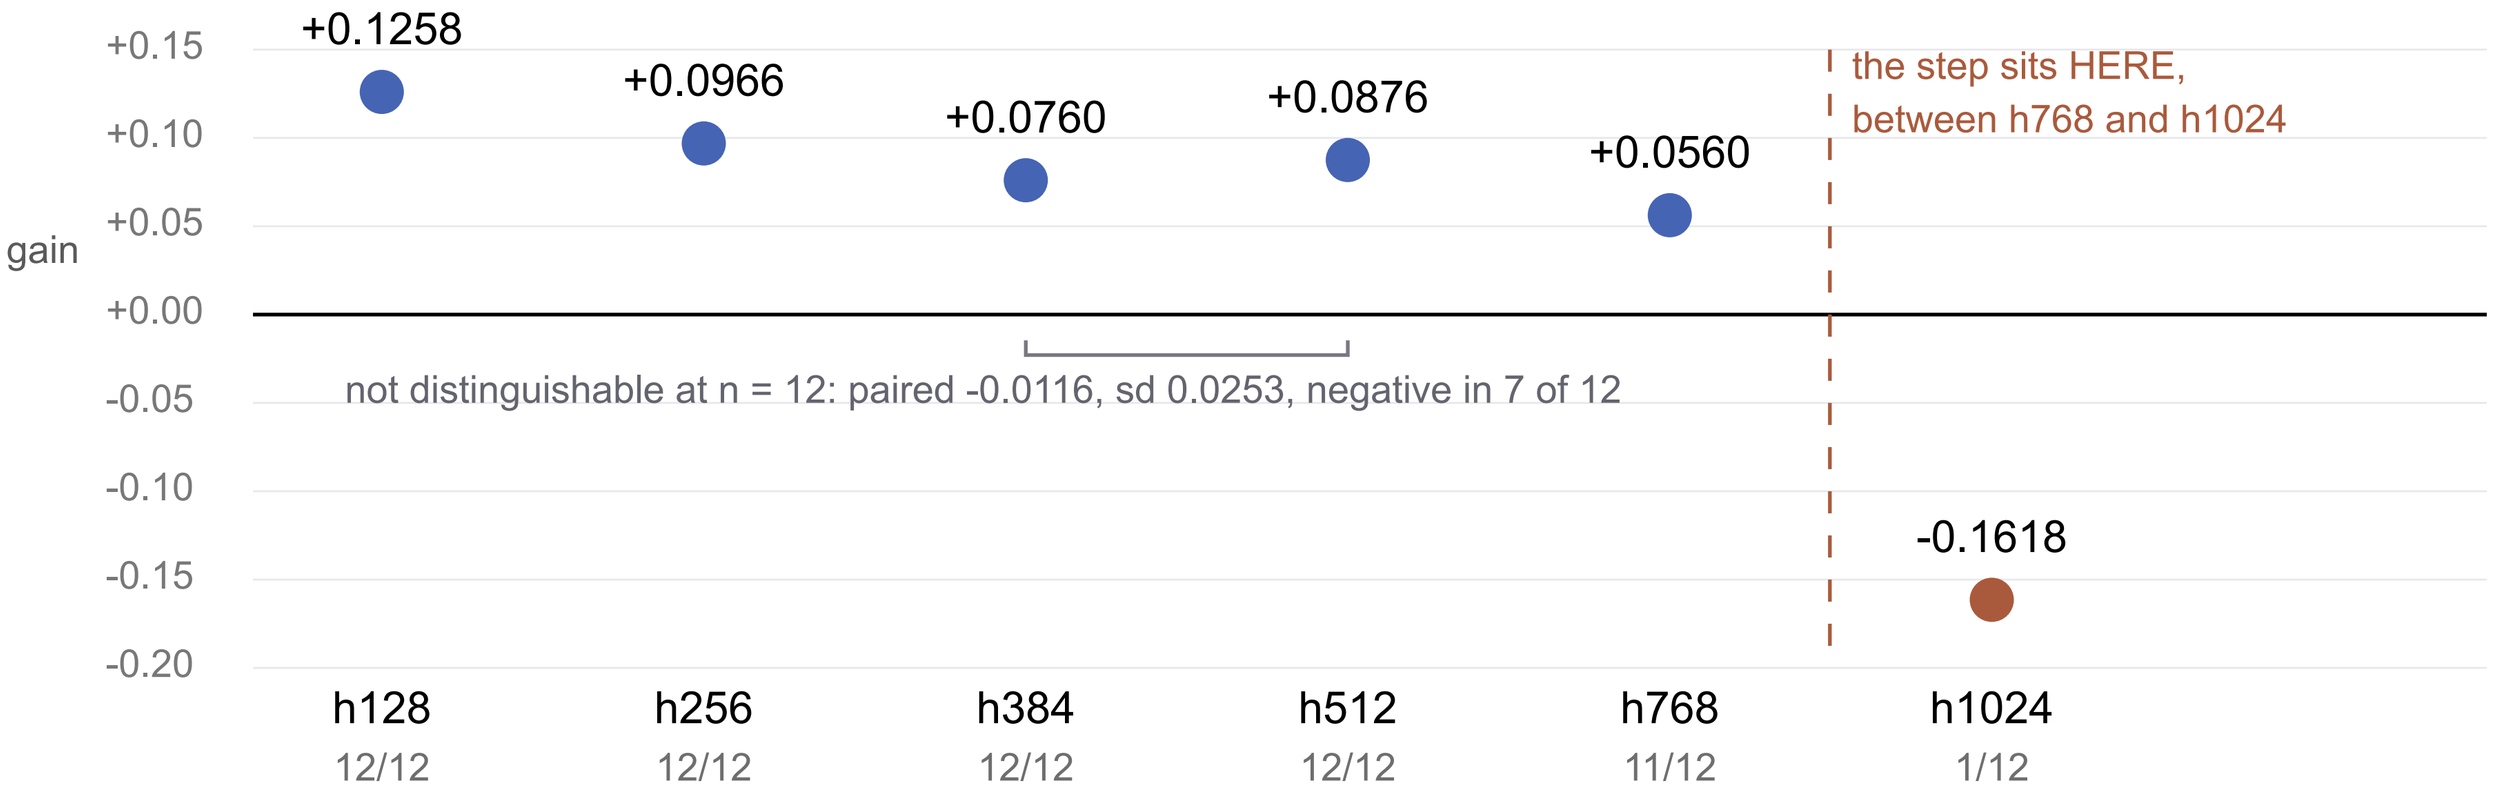

The drop into h1024 is 0.2178, 6.9× the largest gap below it (0.0316), and more than twice the registered 3× bar. That is what makes it a threshold rather than the slope continuing.

THE MECHANISM DOES NOT TRACK THIS CURVE. Spearman rho between these rungs and their difference-in-differences is -0.1429, against a registered bar of +0.829: NOT MET. h768 carries the smallest positive gain on this ladder (+0.0560) and the largest DiD in the campaign (+0.1881). The DiD ladder is Figure 1 Panel D, and is deliberately not superimposed here.

Panel B: located, and not explained. Three preregistered rescue levers at h1024 / d32/L4, n = 12 each.

surrogate scale 0.5	gain -0.2106	positive 0/12	median epoch-mean gradient norm 142.009
surrogate scale 0.25	gain -0.2565	positive 0/12	median epoch-mean gradient norm 151.391
clip-grad-norm 1000.0	gain -0.0904	positive 1/12	median epoch-mean gradient norm 11.660
(the unclipped arm they were to rescue)	gain -0.1618	positive 1/12	median epoch-mean gradient norm 55.494

At h512 the clipping flag is inert: 12/12 cells byte-identical to the archived unclipped cells.

LOCATED BUT UNEXPLAINED. Gradient norms leaving O(1) are a correlate, not a cause. It is a lost fit, not overfitting: the training loss ends 56x above its own minimum.