


Figure 3 — The width ladder, and a threshold rather than a slope

The gain is positive on five rungs and inverts at h1024. n = 12 per rung, d32/L4, e400, seed-paired. Rungs are drawn as points: no curve is fitted and none may be.



The drop into h1024 is 0.2178 — 6.9× the largest gap below it (0.0316), and more than twice the registered 3× bar. That is what makes it a threshold rather than the slope continuing.

No monotonic decay is claimed above the threshold, and no dip at h384: the registration asked for 0.005 separations over quantities inside their own noise floor.

Panel B — located, and not explained. Three preregistered rescue levers at h1024 / d32/L4, n = 12 each.

Every lever is negative, and every one is worse than the arm it was meant to rescue.

surrogate scale 0.5	gain -0.2106	positive 0/12	median epoch-mean gradient norm 142.009
surrogate scale 0.25	gain -0.2565	positive 0/12	median epoch-mean gradient norm 151.391
clip-grad-norm 1000.0	gain -0.0904	positive 1/12	median epoch-mean gradient norm 11.660
(the unclipped arm they were to rescue)	gain -0.1618	positive 1/12	median epoch-mean gradient norm 55.494

Clipping moved the median gradient norm from 55.494 to 11.660 and accuracy did not follow. At h512 the same flag is inert: 12/12 cells byte-identical to the archived unclipped cells.

LOCATED BUT UNEXPLAINED. Gradient norms leaving O(1) are a correlate, not a cause, and overfitting on 8,156 training samples is not excluded by anything in this paper.